

# HADOOP

## 1. Data Ingestion:

- Create a directory in HDFS and transfer the banking dataset from the local system to the HDFS directory.

Solution – Creation of directory

```
C:\hadoop\sbin>hdfs dfs -mkdir /user/Satyam  
C:\hadoop\sbin>hdfs dfs -put C:/dataset/bank.csv /user/Satyam
```

Transfer the banking dataset from the local system to the HDFS directory.

## Browse Directory

/user/Satyam

Go!

Show

25

entries

Search:

<div><input type="checkbox"/></div>	<div><div><div></div><div></div><div></div></div>Permission</div>	<div><div><div></div><div></div><div></div></div>Owner</div>	<div><div><div></div><div></div><div></div></div>Group</div>	<div><div><div></div><div></div><div></div></div>Size</div>	<div><div><div></div><div></div><div></div></div>Last Modified</div>	<div><div><div></div><div></div><div></div></div>Replication</div>	<div><div><div></div><div></div><div></div></div>Block Size</div>	<div><div><div></div><div></div><div></div></div>Name</div>	<div><div><div></div><div></div><div></div></div></div>
<div><input type="checkbox"/></div>	<div><div>-rw-r--r--</div></div>	<div><div>Shri</div></div>	<div><div>supergroup</div></div>	<div><div>366.74 KB</div></div>	<div><div>Oct 04 16:32</div></div>	<div><div>1</div></div>	<div><div>128 MB</div></div>	<div><div>bank.csv</div></div>	<div><div><div></div></div></div>

Showing 1 to 1 of 1 entries

Previous

1

Next

## 2. Data Transformation with MapReduce:

- Write a MapReduce program in Python that calculates the average account balance for each job type.

Solution –

```
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem1/mapper.py /user/Satyam/input1  
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem1/reducer.py /user/Satyam/input1  
C:\hadoop\sbin>hadoop jar C:/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar -files "file:///C:/dataset/Mapreduce/Problem1/mapper.py,file:///C:/dataset/Mapreduce/Problem1/reducer.py" -mapper "python mapper.py" -reducer "python reducer.py" -input /user/Satyam/bank.csv -output /user/Satyam/output1
```

```
C:\hadoop\sbin>hdfs dfs -cat /user/Satyam/output1/part-00000  
admin. 1226.73640167364  
blue-collar 1085.161733615222  
entrepreneur 1645.125  
housemaid 2083.8035714285716  
management 1766.9287925696594  
retired 2319.191304347826  
self-employed 1392.4098360655737  
services 1103.9568345323742  
student 1543.8214285714287  
technician 1330.99609375  
unemployed 1089.421875  
unknown 1501.7105263157894
```

Name – Satyam Kulkarni

- Write another MapReduce program that counts the number of individuals with and without a housing loan in each education category.

```
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem2/mapper.py /user/Satyam/input2
put: `/user/Satyam/input2/mapper.py': File exists

C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem2/reducer.py /user/Satyam/input2

C:\hadoop\sbin>hadoop jar C:/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar -files "file:///C:/dataset/Mapreduce/Problem2/mapper.py,file:///C:/dataset/Mapreduce/Problem2/reducer.py" -mapper "python mapper.py" -reducer "python reducer.py" -input /user/Satyam/bank.csv -output /user/Satyam/output2
```

```
C:\hadoop\sbin>hdfs dfs -cat /user/Satyam/output2/part-00000
primary 94      583
secondary 416    1889
tertiary 173    1176
unknown 7      179
```

- Perform a MapReduce job to determine the number of clients contacted in each month and their subscription status to term deposits ('y' column).

```
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem3/mapper.py /user/Satyam/input3

C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem3/reducer.py /user/Satyam/input3

C:\hadoop\sbin>hadoop jar C:/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar -files "file:///C:/dataset/Mapreduce/Problem3/mapper.py,file:///C:/dataset/Mapreduce/Problem3/reducer.py" -mapper "python mapper.py" -reducer "python reducer.py" -input /user/Satyam/bank.csv -output /user/Satyam/output3
```

```
C:\hadoop\sbin>hdfs dfs -cat /user/Satyam/output3/part-00000
apr      56      236
aug      79      553
dec       8       11
feb      38      183
jan      16      131
jul      61      644
jun      55      475
mar      21       27
may      93     1304
nov      38      350
oct      37       42
sep      16       35
```

### 3. Data Analysis with MapReduce:

- Analyze the average duration of contact (in seconds) per campaign outcome ('poutcome').

```
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem4/mapper.py /user/Satyam/input4

C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem4/reducer.py /user/Satyam/input4

C:\hadoop\sbin>hadoop jar C:/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar -files "file:///C:/dataset/Mapreduce/Problem4/mapper.py,file:///C:/dataset/Mapreduce/Problem4/reducer.py" -mapper "python mapper.py" -reducer "python reducer.py" -input /user/Satyam/bank.csv -output /user/Satyam/output4
```

Name – Satyam Kulkarni

```
C:\hadoop\sbin>hdfs dfs -cat /user/Satyam/output4/part-00000
failure 254.38367346938776
other 273.83248730964465
success 338.6356589147287
unknown 262.1031039136302
```

- Examine the relationship between the age of clients and their balance, and present findings in a summarized form.

```
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem5/mapper.py /user/Satyam/input5
C:\hadoop\sbin>hdfs dfs -put C:/dataset/Mapreduce/Problem5/reducer.py /user/Satyam/input5
```

```
C:\hadoop\sbin>hadoop jar C:/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar -files "file:///C:/dataset/Mapreduce/Problem5/mapper.py,file:///C:/dataset/Mapreduce/Problem5/reducer.py" -mapper "python mapper.py" -reducer "python reducer.py" -input /user/Satyam/bank.csv -output /user/Satyam/output5
```

```
C:\hadoop\sbin>hdfs dfs -cat /user/Satyam/output5/part-00000
19      393.5
20      661.3333333333334
21      1774.2857142857142
22      1455.3333333333333
23      2117.95
24      634.625
25      1240.0681818181818
26      788.5584415584416
27      851.7765957446809
28      1025.0970873786407
29      1261.8762886597938
30      1113.0333333333333
31      1288.4824120603016
32      1256.549107142857
33      1545.4139784946237
34      1111.5367965367966
35      1192.8277777777778
36      1226.8936170212767
37      1463.9192546583852
38      1718.993710691824
39      1104.8615384615384
40      1399.5070422535211
41      1505.7925925925927
42      1612.3617021276596
43      1807.8347826086956
44      1836.5523809523809
45      1187.3660714285713
46      998.7731092436975
47      1363.0462962962963
48      1462.359649122807
49      1591.107142857143
50      1645.0594059405942
51      1528.5714285714287
52      782.2906976744187
53      1588.3085106382978
54      1656.661971830986
55      1244.9444444444443
56      2120.135135135135
```

## HIVE

### 1. Data Ingestion and Table Creation:

- Create a Hive database named **banking\_data**.
- Define and create a Hive table **client\_info** with appropriate data types for the **bank.csv** dataset.

```
hive> CREATE TABLE client_info(  
  > age INT,  
  > job STRING,  
  > marital STRING,  
  > education STRING,  
  > default STRING,  
  > balance INT,  
  > housing STRING,  
  > loan STRING,  
  > contact STRING,  
  > day INT,  
  > month STRING,  
  > duration INT,  
  > campaign INT,  
  > pdays INT,  
  > previous INT,  
  > poutcome STRING,  
  > y STRING  
  > )  
  > ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ','  
  > STORED AS TEXTFILE;  
2024-10-05T15:41:42,467 INFO [mai  
40dcc72b-e2d0-48e9-97b4-9f1ef28cd  
2024-10-05T15:41:42,469 INFO [mai  
0-48e9-97b4-9f1ef28cdc23 main  
OK
```

- Load the data from the **bank.csv** file into the **client\_info** table.

```
hive> LOAD DATA INPATH '/user/Satyam/bank.csv' INTO TABLE client_info;  
2024-10-05T15:58:30,054 INFO [main] org.apache.hadoop.hive.conf.HiveConf  
40dcc72b-e2d0-48e9-97b4-9f1ef28cdc23  
2024-10-05T15:58:30,060 INFO [main] org.apache.hadoop.hive ql.session.Ses  
0-48e9-97b4-9f1ef28cdc23 main  
Loading data to table banking_data.client_info  
OK
```

```
hive> SELECT * FROM client_info LIMIT 10;
```

Name – Satyam Kulkarni

id	job	marital	education	default	NULL	housing	loan	contact	NULL	month	NULL	NULL	NULL	N
30	unemployed	married	primary	no	1787	no	no	cellular		19	oct	79	1	-
33	services	married	secondary	no	4789	yes	yes	cellular		11	may	220	1	
35	management	single	tertiary	no	1350	yes	no	cellular		16	apr	185	1	
30	management	married	tertiary	no	1476	yes	yes	unknown	3	jun	199	4	-	
59	blue-collar	married	secondary	no	0	yes	no	unknown	5	may	226	1	-	
35	management	single	tertiary	no	747	no	no	cellular		23	feb	141	2	
36	self-employed	married	tertiary	no	307	yes	no	cellular		14	may	341	1	
39	technician	married	secondary	no	147	yes	no	cellular		6	may	151	2	
41	entrepreneur	married	tertiary	no	221	yes	no	unknown	14	may	57	2	-	

Time taken: 6.784 seconds, Fetched: 10 row(s)

## 2. Basic Data Exploration:

- Write a HiveQL query to count the total number of clients in the dataset.

Ans:

```
hive> select count(*) as total_clients from client_info;
```

```
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU
Total MapReduce CPU Time Spent: 12 seconds 246 ms
OK
4522
Time taken: 107.457 seconds, Fetched: 1 row(s)
```

- Display the first 10 rows of the dataset.

Ans:

id	job	marital	education	default	NULL	housing	loan	contact	NULL	month	NULL	NULL	NULL	N
30	unemployed	married	primary	no	1787	no	no	cellular		19	oct	79	1	-
33	services	married	secondary	no	4789	yes	yes	cellular		11	may	220	1	
35	management	single	tertiary	no	1350	yes	no	cellular		16	apr	185	1	
30	management	married	tertiary	no	1476	yes	yes	unknown	3	jun	199	4	-	
59	blue-collar	married	secondary	no	0	yes	no	unknown	5	may	226	1	-	
35	management	single	tertiary	no	747	no	no	cellular		23	feb	141	2	
36	self-employed	married	tertiary	no	307	yes	no	cellular		14	may	341	1	
39	technician	married	secondary	no	147	yes	no	cellular		6	may	151	2	
41	entrepreneur	married	tertiary	no	221	yes	no	unknown	14	may	57	2	-	

Time taken: 6.784 seconds, Fetched: 10 row(s)

## 3. Data Filtering and Sorting:

- Retrieve all records of clients who are married and have a personal loan.

Ans:

```
hive> select * from client_info where marital = 'married' and loan = 'yes';
```

Name – Satyam Kulkarni

```
OK
33      services      married secondary      no      4789      yes      yes      cellular      11      may      220      1      33
39      4      failure no      management      married tertiary      no      1476      yes      yes      unknown 3      jun      199      4      -1      0
30      unknown no      services      married primary no      -88      yes      yes      cellular      17      apr      313      1      147      2
43      failure no      blue-collar      married secondary      no      360      yes      yes      cellular      29      jan      89      1      24
31      1      failure no      management      married tertiary      no      194      no      yes      cellular      29      aug      189      2      -1
40      0      unknown no      self-employed      married secondary      no      784      no      yes      cellular      30      jul      149      2      -1
56      0      unknown no      admin.      married secondary      no      105      no      yes      cellular      21      aug      74      2      -1      0
53      unknown no      management      married secondary      no      82      no      yes      telephone      4      feb      140      1      -1
57      0      unknown no      blue-collar      married primary no      -516      no      yes      telephone      8      jul      554      3      -1      0
41      unknown no      management      married secondary      no      0      no      yes      cellular      7      jul      630      3      -1
41      0      unknown no      blue-collar      married secondary      no      427      yes      yes      unknown 9      jun      371      3      -1      0
37      unknown no      self-employed      married secondary      no      217      yes      yes      cellular      15      jul      317      5      -1
32      0      unknown no      blue-collar      married secondary      no      -231      no      yes      cellular      15      jul      779      2      -1
42      0      admin.      married secondary      no      323      yes      yes      unknown 8      may      280      2      -1      0      un
known no
35      0      management      married tertiary      no      106      no      yes      cellular      11      aug      588      2      -1
56      0      unknown no      retired married primary no      1906      no      yes      unknown 19      jun      45      9      -1      0      unknown no
43      services      married secondary      no      978      yes      yes      yes      unknown 26      may      82      2      -1      0
unknown no
43      admin.      married secondary      no      -465      yes      yes      cellular      23      jul      166      1      -1      0
unknown no
```

```
Time taken: 1.168 seconds, Fetched: 453 row(s)
```

- List the top 10 clients with the highest balance, displaying their job, marital status, and balance.

Ans:

```
hive> select job,marital,balance from client_info order by balance desc limit 10;
```

```
OK
retired married 71188
entrepreneur      married 42045
technician        single  27733
management        married 27359
technician        married 27069
housemaid         single  26965
retired married 26452
services          married 26394
management        divorced      26306
retired single 25824
Time taken: 85.493 seconds, Fetched: 10 row(s)
```

#### 4. Data Aggregation and Grouping:

- Calculate the average age of clients for each job category.

Ans:

```
hive> select job, avg(age) as average_age from client_info group by job;
```

```
OK
admin. 39.68200836820084
blue-collar 40.15644820295983
entrepreneur 42.01190476190476
housemaid 47.339285714285715
job NULL
management 40.54076367389061
retired 61.869565217391305
self-employed 41.45355191256831
services 38.57074340527578
student 26.821428571428573
technician 39.470052083333336
unemployed 40.90625
unknown 48.10526315789474
Time taken: 123.358 seconds, Fetched: 13 row(s)
```

- Find the total number of clients for each education level who have defaulted on credit.

Ans:

```
hive> select education, default, count(*) as total_defaulted_client from client_info where default = 'yes' group by education, default
```

```
OK
primary yes 10
secondary yes 46
tertiary yes 17
unknown yes 3
Time taken: 95.775 seconds, Fetched: 4 row(s)
```

## 5. Complex Queries for Insights:

- Identify the top 5 job categories with the highest average balance and the percentage of clients in each of these job categories who have subscribed to a term deposit.

Ans:

We directly calculate the average balance (AVG(ci.balance)) for each job category.

We also calculate the total number of clients (COUNT(\*)) and the number of clients who have subscribed to a term deposit (SUM(CASE WHEN ci.y = 'yes' THEN 1 ELSE 0 END)).

We calculate the subscription rate as the percentage of subscribed clients out of the total number of clients.

The results are grouped by job category and sorted in descending order of average balance.

We limit the output to the top 5 job categories with the highest average balance using the LIMIT clause.

This query will give you the top 5 job categories with the highest average balance and the percentage of clients in each of these job categories who have subscribed to a term deposit.

```
hive> select s.job, s.avg_balance, (s.subscribed_clients/s.total_clients)*100 as subscription_percentage from (select job, avg(balance) as avg_balance, count(*) as total_clients, sum(case when y = 'yes' then 1 else 0 end) as subscribed_clients from client_info group by job order by avg_balance desc limit 5) s;
```



```
OK
retired 2319.191304347826      23.47826086956522
housemaid 2083.8035714285716    12.5
management 1766.9287925696594  13.519091847265221
entrepreneur 1645.125          8.928571428571429
student 1543.8214285714287     22.61904761904762
Time taken: 172.02 seconds, Fetched: 5 row(s)
```

- Determine the month with the highest number of contacts and the success rate of the campaign in that month (percentage of clients who subscribed to a term deposit).

Ans:

We directly calculate the success rate for each month by counting the total number of contacts and the number of clients who subscribed to a term deposit within each month.

We use the GROUP BY clause to group the data by month.

The results are sorted in descending order of the number of contacts, and we limit the output to only the first row, which represents the month with the highest number of contacts.

This query will give you the month with the highest number of contacts and the success rate of the campaign in that month.

```
hive> select month, total_contacts, (successful_contacts/total_contacts)*100 as success_rate from (select month, count(*) as total_contacts, sum(case when y = 'yes' then 1 else 0 end) as successful_contacts from client_info group by month order by total_contacts desc limit 1) as top_month;
```

```
OK
may      1398      6.652360515021459
Time taken: 172.717 seconds, Fetched: 1 row(s)
```

## 6. Correlation Analysis:

- Calculate the correlation between age and balance for the clients.

Ans:

In this simplified query:

1. We calculate the numerator of the correlation formula: the sum of the products of age and balance minus the count of records times the average of age times the average of balance.
2. We calculate the denominator of the correlation formula: the square root of the difference between the sum of the squares of age and the count times the square of the average of age, multiplied by the square root of the difference between the sum of the squares of balance and the count times the square of the average of balance.
3. We divide the numerator by the denominator to get the correlation coefficient.

```
hive> select corr(age, balance) as age_balance_correlation from client_info;
```

```
OK
0.08382014224477742
Time taken: 94.424 seconds, Fetched: 1 row(s)
```



## 7. Trend Analysis:

- Analyze the year-over-year trend in the number of clients contacted.

Ans:

In this query:

We use the SUBSTRING function to extract the year from the month column. Assuming the month column is in the format "YYYY-MM", we extract the first four characters to get the year.

We count the number of clients contacted (COUNT(\*)) for each year.

We group the results by the extracted year.

Finally, we order the results by year to see the trend over time.

This query will give you the year-over-year trend in the number of clients contacted based on the data in your client\_info table.

```
hive> select substring(month,1,4) as year, count(*) as num_clients_contacted from client_info group by substring(month,1,4) order by year;
```

```
OK
apr      293
aug      633
dec       20
feb      222
jan      148
jul      706
jun      531
mar       49
may     1398
mont       1
nov      389
oct       80
sep       52
Time taken: 163.069 seconds, Fetched: 13 row(s)
```

## 8. Anomaly Detection:

- Identify any unusual patterns in the average yearly balance across different education levels.

Ans:

We calculate the average yearly balance for each education level.

We also calculate the overall average yearly balance for each year using the window function AVG(AVG(balance)) OVER (PARTITION BY SUBSTRING(month, 1, 4)).

We calculate the z-score for each average yearly balance within each education level by subtracting the overall average yearly balance and dividing by the standard deviation, both calculated over the partition of years.

By examining the z-scores, you can identify any unusual patterns in the average yearly balance across different education levels. A z-score significantly higher or lower than zero indicates that the average yearly balance for a particular education level is unusually high or low compared to the overall average yearly balance.

```
hive> select year, education, (avg_yearly_balance - overall_avg_balance) / stddev_balance as z_score from (select substring(month,1,4) as year, education, avg(balance) as avg_yearly_balance, avg(avg(balance)) over (partition by substring(month,1,4)) as overall_avg_balance, STDDEV(avg(balance)) over (partition by substring(month,1,4)) as stddev_balance from client_info group by substring(month,1,4), education) as subquery;
```

Name – Satyam Kulkarni

```
OK
apr      primary 1.267794378462834
apr      secondary -0.7235770026264813
apr      tertiary 0.6555354947097891
apr      unknown -1.1997528705461413
aug      primary -0.8113812295586758
aug      secondary -0.5775732890112949
aug      tertiary -0.31633238372712924
aug      unknown 1.7052869022970998
dec      primary -0.8088635843923103
dec      secondary -0.397568378390461
dec      tertiary -0.5058241544774784
dec      unknown 1.7122561172602497
feb      primary -1.0701923004119014
feb      secondary -0.9216091477916096
feb      tertiary 0.8917649610011248
feb      unknown 1.1000364872023862
jan      primary 1.1404156384061253
jan      secondary -0.4369029752604172
jan      tertiary 0.7115176049750513
jan      unknown -1.41503026812076
jul      primary 0.9649291030842655
jul      secondary -0.4258456259680814
jul      tertiary 0.9014112560610348
jul      unknown -1.4404947331772213
jun      primary -0.15682779176634196
jun      secondary -1.0365919570822748
jun      tertiary 1.6428352994501707
jun      unknown -0.44941555060155386
mar      primary -1.0722876073529977
mar      secondary -0.2462824810074777
mar      tertiary 1.6391398006892093
mar      unknown -0.3205697123287344
may      primary -1.0048248682462826
may      secondary -0.9457054758622359
may      tertiary 0.6640708393343474
may      unknown 1.2864595047741691
```

```
mont     education NULL
nov      primary -0.8137802862217836
nov      secondary 0.2770762379753206
nov      tertiary 1.5167439470393023
nov      unknown -0.9800398987928378
oct      primary 1.7097889755337383
oct      secondary -0.37052787477325194
oct      tertiary -0.5238214197432853
oct      unknown -0.8154396810172013
sep      primary -1.3086820436917062
sep      secondary -0.4707154583719365
sep      tertiary 0.3984500713422833
sep      unknown 1.38094743072136
Time taken: 94.281 seconds, Fetched: 49 row(s)
```

### 9. Advanced Analysis:

- Analyze the impact of previous campaign outcomes (**poutcome**) on the current campaign's success. Calculate the subscription rate (to term deposits) for each **poutcome** category.

Ans:

In this query:

1. We group the data by the poutcome column to analyze the impact of previous campaign outcomes.
2. We count the total number of clients (total\_clients) and the number of clients who subscribed to term deposits (subscribed\_clients) for each poutcome category.
3. We calculate the subscription rate (subscription\_rate) as the percentage of clients who subscribed to term deposits out of the total number of clients for each poutcome category.

```
hive> select poutcome, count(*) as total_clients, sum(case when y = 'yes' then 1 else 0 end) as subscribed_clients, round(sum(case when y = 'yes' then 1 else 0 end) * 100.0/count(*),2) as subscription_rate from client_info group by poutcome order by subscription_rate desc;
```

```
OK
success 129      83      64.34
other   197      38      19.29
failure 490      63      12.86
unknown 3705     337      9.10
poutcome 1         0       0.00
Time taken: 170.274 seconds, Fetched: 5 row(s)
```

- Compare the average contact duration for clients who subscribed and who did not subscribe to a term deposit.

Ans:

In this query:

1. We group the data by the subscription status (y column), which indicates whether the client subscribed to a term deposit.
2. We calculate the average contact duration (avg\_contact\_duration) for each group separately.

This query will provide you with the average contact duration for clients who

**Name – Satyam Kulkarni**

subscribed and who did not subscribe to a term deposit, allowing you to compare the

contact durations between the two groups.

```
hive> select y as subscription_status, avg(duration) as avg_contact_duration from client_info group by y;
```

```
OK
no      226.3475
y       NULL
yes     552.7428023032629
Time taken: 87.218 seconds, Fetched: 3 row(s)
```

### **Submission Guidelines:**

- Make a copy of this doc file.
- Perform the analysis in your local system using Hadoop and Hive and provide screenshots of both the **code** and the **output** under each question.
- Upload the doc file with other files and submit it in the submission dashboard.